

- Divide-and-conquer

- ▶ Each *task* recursively creates n tasks that divide the problem into subproblems
- ▶ Each task t then waits for all n tasks to finish and then may 'combine' the responses
- ▶ At some point the recursion stops (at the bottom of the tree), and some sequential kernel is executed
- ▶ Then the result is propagated upward in the tree recursively
- ▶ Examples: fibonacci, quick sort, ...

Task Parallelism

- State-space search
 - ▶ Each *task* recursively creates n tasks to partition the search space
 - ▶ If the problem is one-solution search, as soon as a task encounters a solution, the program may need to terminate
 - ★ Kill-chasing problem
- All-solution search may require behaviour much like divide-and-conquer where values are combined
 - ▶ Example: all-solution n queens
 - ▶ Number of solutions are accumulated recursively up the tree

Fibonacci Example

- Each `Fib` chore is a task that performs one of two actions:
 - ▶ Creates two new `Fib` chores to compute $fib(n - 1)$ and $fib(n - 2)$ and then waits for the response, adding up the two responses when they arrive
 - ★ After both arrive, sends a response message with the result to the parent task
 - ★ Or prints the value and calls `CkExit()` if it is the root
 - ▶ If $n = 1$ or $n = 0$ (passed down from the parent) it sends a response message with n back to the parent task

Fibonacci Example

```
mainmodule fib {  
  mainchare Main {  
    entry Main(CkArgMsg* m);  
  };  
  
  chare Fib {  
    entry Fib(int n, bool isRoot, CProxy_Fib parent);  
    entry void response(int value);  
  };  
};
```

Fibonacci Example

```
struct Main : public CBase_Main {
    Main(CkArgMsg* m) {
        CProxy_Fib::ckNew(atoi(m->argv[1]), true, CProxy_Fib());
    }
};

struct Fib : public CBase_Fib {
    CProxy_Fib parent; bool isRoot; int result, count;

    Fib(int n, bool isRoot_, CProxy_Fib parent_)
        : parent(parent_), isRoot(isRoot_), result(0), count(n < 2 ? 1 : 2) {


        if (n < 2) response(n);
        else {
            CProxy_Fib::ckNew(n - 1, false, thisProxy);
            CProxy_Fib::ckNew(n - 2, false, thisProxy);
        }
    }

    void response(int val) {
        result += val;
        if (--count == 0) {
            if (isRoot) {
                CkPrintf("Fibonacci number is: %d\n", result);
                CkExit();
            } else {
                parent.response(result);
                delete this;
            }
        }
    }
};
```


Fibonacci Execution

fib(5)


Fibonacci Execution



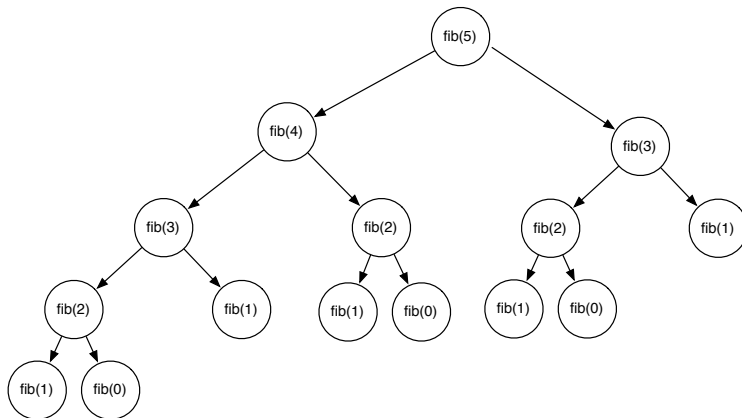
Fibonacci Execution



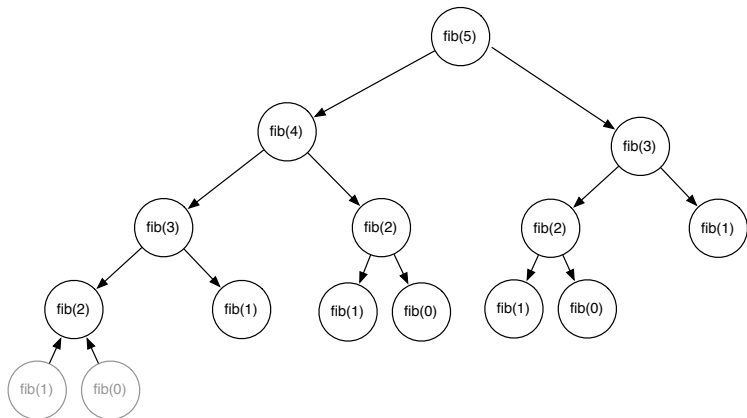
Fibonacci Execution




Fibonacci Execution




Fibonacci Execution




Fibonacci Execution



Fibonacci Execution



Fibonacci Execution



Fibonacci Execution


fib(5)

Fibonacci Performance

- How much work/computation does each chare do in this example?
- What are some of the overheads of this approach?
- Is there way we can reduce/amortize the overhead?

Possible Solution

- Set a sequential threshold in the computational tree
 - ▶ Past this threshold (i.e. when $n < threshold$), instead of constructing two new chares, compute the fibonacci sequentially



- $fib(5), fib(4)$ are fine grains, $fib(3), fib(2)$ are coarser grains
- The coarser grains now amortize the cost of the fine-grained execution

Fibonacci w/Threshold Example

```
#define THRESHOLD 10

struct Main : public CBase_Main { /* ... same as before ... */};

struct Fib : public CBase_Fib {
    CProxy_Fib parent; bool isRoot; int result, count;

    Fib(int n, bool isRoot_, CProxy_Fib parent_)
        : parent(parent_), isRoot(isRoot_), result(0), count(n < THRESHOLD ? 1 : 2) {

        if (n < THRESHOLD) response(seqFib(n));
        else {
            CProxy_Fib::ckNew(n - 1, false, thisProxy);
            CProxy_Fib::ckNew(n - 2, false, thisProxy);
        }
    }

    int seqFib(int n) { return (n < 2) ? n : seqFib(n - 1) + seqFib(n - 2); }

    void response(int val) {
        result += val;
        if (--count == 0) {
            if (isRoot) {
                CkPrintf("Fibonacci number is: %d\n", result);
                CkExit();
            } else {
                parent.response(result);
                delete this;
            }
        }
    }
};
```

Amdahl's Law and Grainsize

- Original “law”:

- ▶ If a program has $K\%$ sequential section, then speedup is limited to $\frac{100}{K}$.
- ★ If the rest of the program is parallelized completely

- Grainsize corollary:


- ▶ If any individual piece of work is $> K$ time units, and the sequential program takes T_{seq} ,
- ★ Speedup is limited to $\frac{T_{seq}}{K}$

- So:

- ▶ Examine performance data via histograms to find the sizes of remappable work units
- ▶ If some are too big, change the decomposition method to make smaller units

Grainsize

- (working) Definition: the amount of computation per potentially parallel event (task creation, enqueue/dequeue, messaging, locking. .)



Grainsize and Overhead

- What is the ideal grainsize?
- Should it depend on the number of processors?

$$T_1 = T \left(1 + \frac{v}{g} \right)$$

$$T_p = \max \left\{ g, \frac{T_1}{p} \right\}$$


$$T_p = \max \left\{ g, \frac{T \left(1 + \frac{v}{g} \right)}{p} \right\}$$

v : overhead per message,

T_p : p processor completion time

g : grainsize (computation per message)

Grainsize and Scalability



Rules of thumb for grainsize

- Make it as small as possible, as long as it amortizes the overhead
- More specifically, ensure:
 - ▶ Average grainsize is greater than kv (say $10v$)
 - ▶ No single grain should be allowed to be too large
 - ★ Must be smaller than $\frac{T}{p}$, but actually we can express it as:
 - ★ Must be smaller than kmv (say $100v$)
- Important corollary:
 - ▶ You can be at close to optimal grainsize without having to think about p , the number of processors

How to determine/ensure grainsize

- Compiler techniques can help, but only in some cases
 - ▶ Note that they don't need precise determination of grainsize, just one that will satisfy a broad inequality
 - ★ $kv < g < mkv$ ($10v < g < 100v$)