


Observations: Exascale applications

- Development of new models must be driven by the needs of exascale applications
 - Multi-resolution
 - Multi-module (multi-physics)
 - Dynamic/adaptive: to handle application variation
 - Adapt to a volatile computational environment
 - Exploit heterogeneous architecture
 - Deal with thermal and energy considerations
- So? Consequences:
 - Must support automated resource management
 - Must support interoperability and parallel composition


The Execution Model

- Several objects live on a single “processor”
 - We will come back what we mean by a processor.
 - For now, think of it as a core
 - As a result,
 - the method invocations directed at objects on that processor will have to be stored in a pool,
 - And a user-level scheduler will select one invocation from the Quque and and runs it completion





Message-driven Execution

- Execution is triggered by availability of a “message” (a method invocation)
- When an entry method executes,
 - it may generate messages for other objects
 - The RTS deposits them in the message Q on the target processor



Utility for Multi-cores, Many-cores, Accelerators:

- Objects connote and promote locality
- Message-driven execution is
 - A strong principle of prediction for data and code use
 - Much stronger than principle of locality
 - Can use to scale memory wall:
 - Prefetching of needed data:
 - into scratch pad memories, for example





Impact on communication

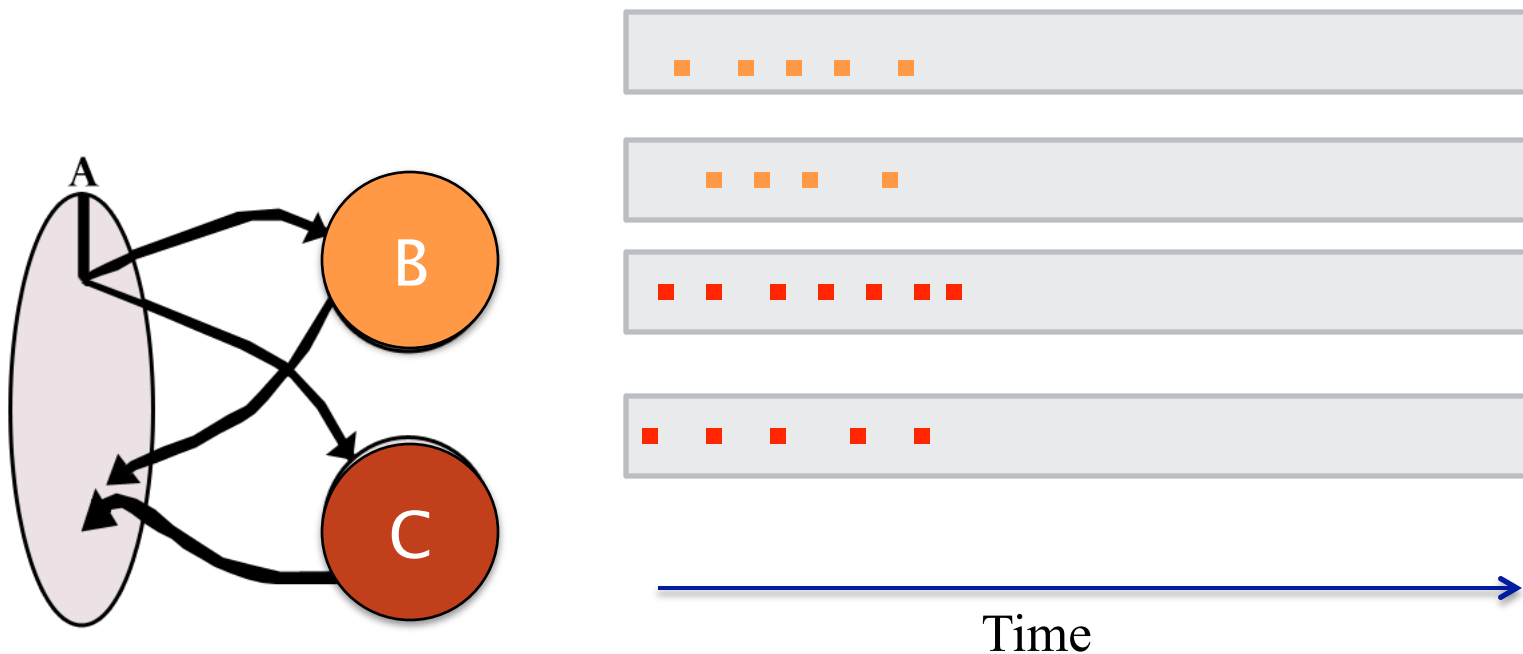
- Current use of communication network:
 - Compute–communicate cycles in typical MPI apps
 - So, the network is used for a fraction of time,
 - and is on the critical path
- So, current *communication networks are over-engineered for by necessity*
- With overdecomposition
 - Communication is spread over an iteration
 - Also, adaptive overlap of communication and computation

Example: Stencil Computation



- Consider a simple stencil computation
 - With traditional design based on traditional methods (e.g. MPI-based)
 - Each processor has a chunk, which alternates between computing and communicating
 - With Charm++
 - Multiple chunks on each processor
 - Wait time for each chunk overlapped with useful computation for the other
 - Communication spreads over time
 - A Simple schematic timeline for both approaches?





Without message-driven execution
(and virtualization), you get either:
Space-division



OR: Sequentialization



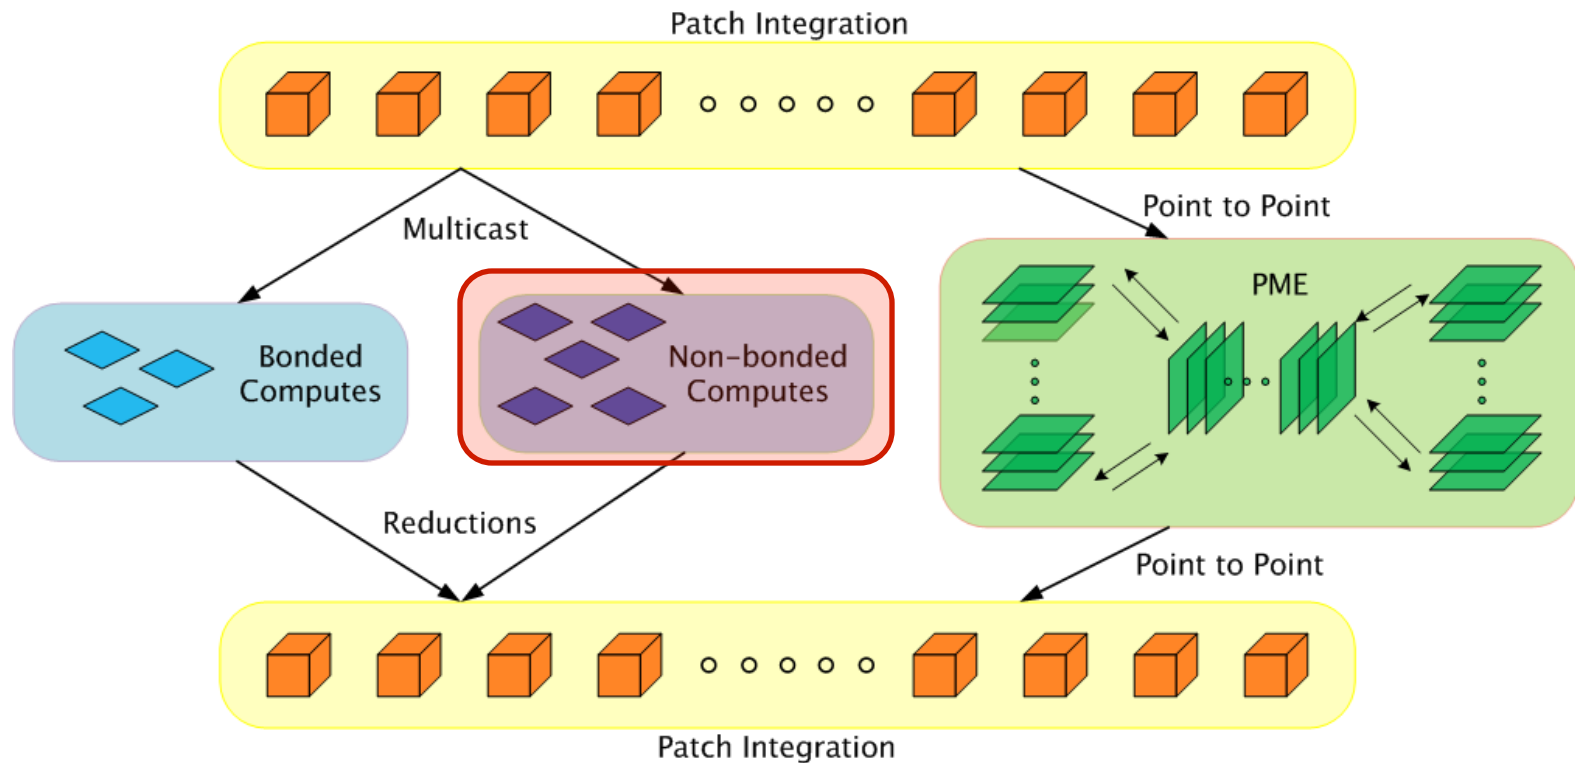
Parallel Composition: $A1; (B \parallel C); A2$



Recall: Different modules, written in different languages/paradigms, can overlap in time and on processors, without programmer having to worry about this explicitly

MD Parallelization Using Charm++

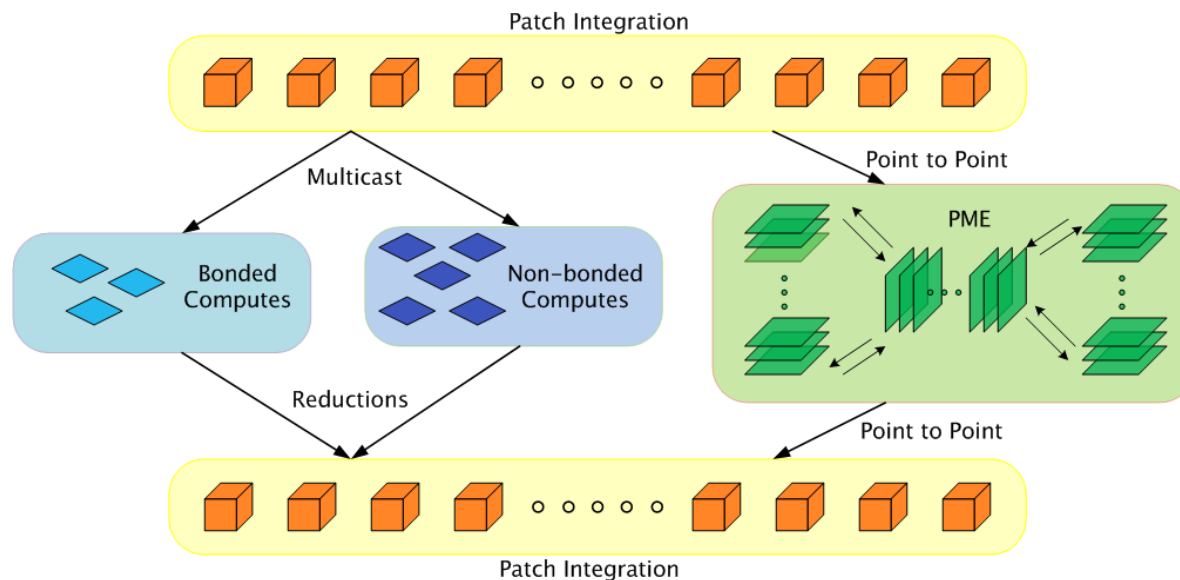
The computation is decomposed into “natural” objects of the application, which are assigned to processors by Charm++ RTS

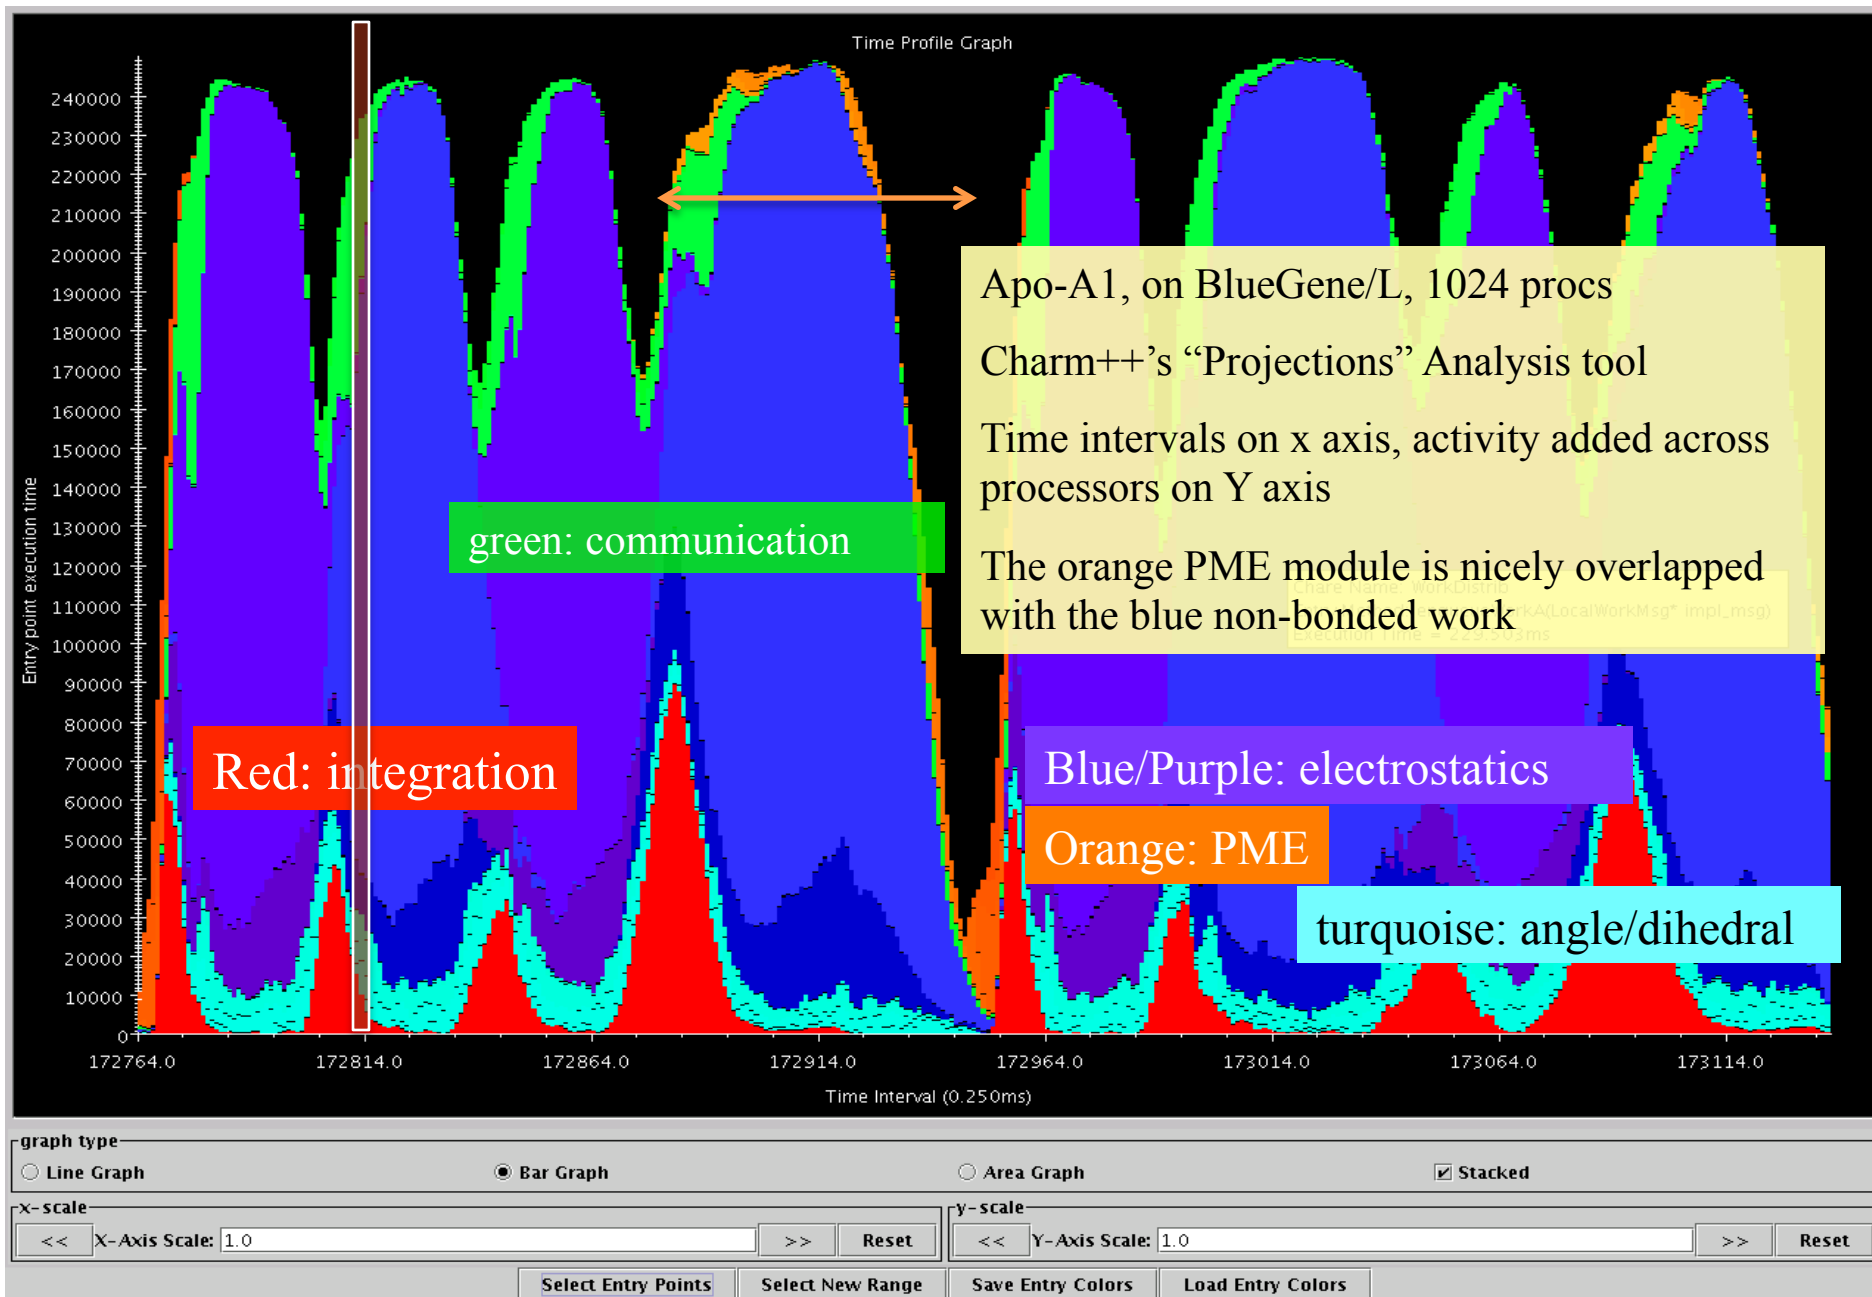


MD Parallelization Using Charm++

The PME module is communication-dominated, with a long critical path, and small computation'

The non-bonded force computation module has a lot of floating point work, but relatively small computation








Time

Decomposition Independent of numCores

- Rocket simulation example under traditional MPI



- With migratable-objects:



- Benefit: load balance, communication optimizations, modularity

Migratability

- Once the programmer has written the code without reference to processors
 - With all the communication expressed as that between objects
- The system is free to migrate the objects across processors as and when it pleases
 - It must ensure it can deliver method invocations to the objects, wherever they go
 - This migratability turns out to be a key attribute for empowering an adaptive runtime system