

Programming

Cell & Larrabee

Using Accelerated Entry Methods

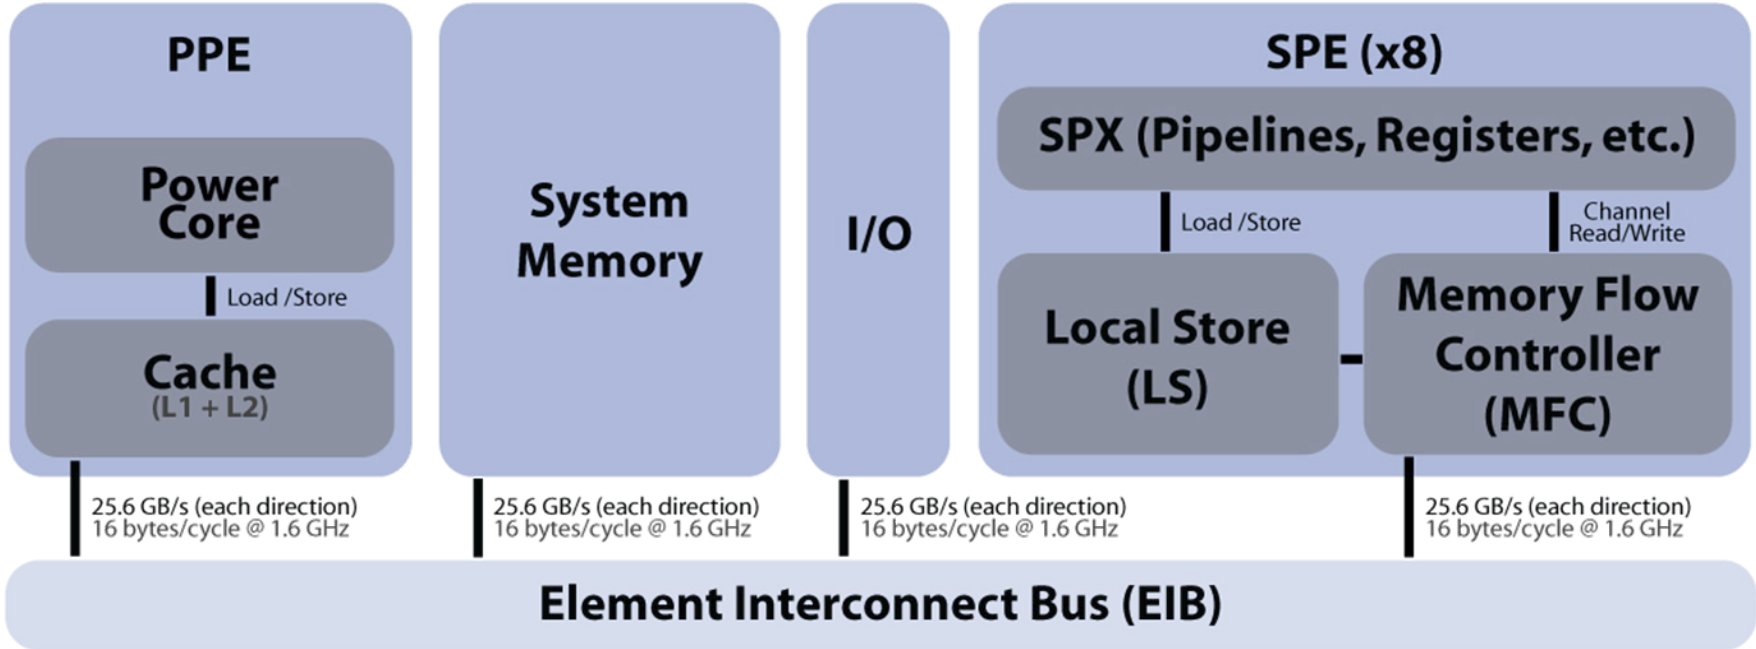
David Kunzman
Parallel Programming Lab
University of Illinois at Urbana-Champaign
11.30.2009

Motivation/Goals

- Allow Charm++ programs to use accelerators
 - Currently targeting Cell and Larrabee
 - Difficult to program
- Ease programmer burden
 - Single programming model for all processing cores
 - Portability
 - Modularity



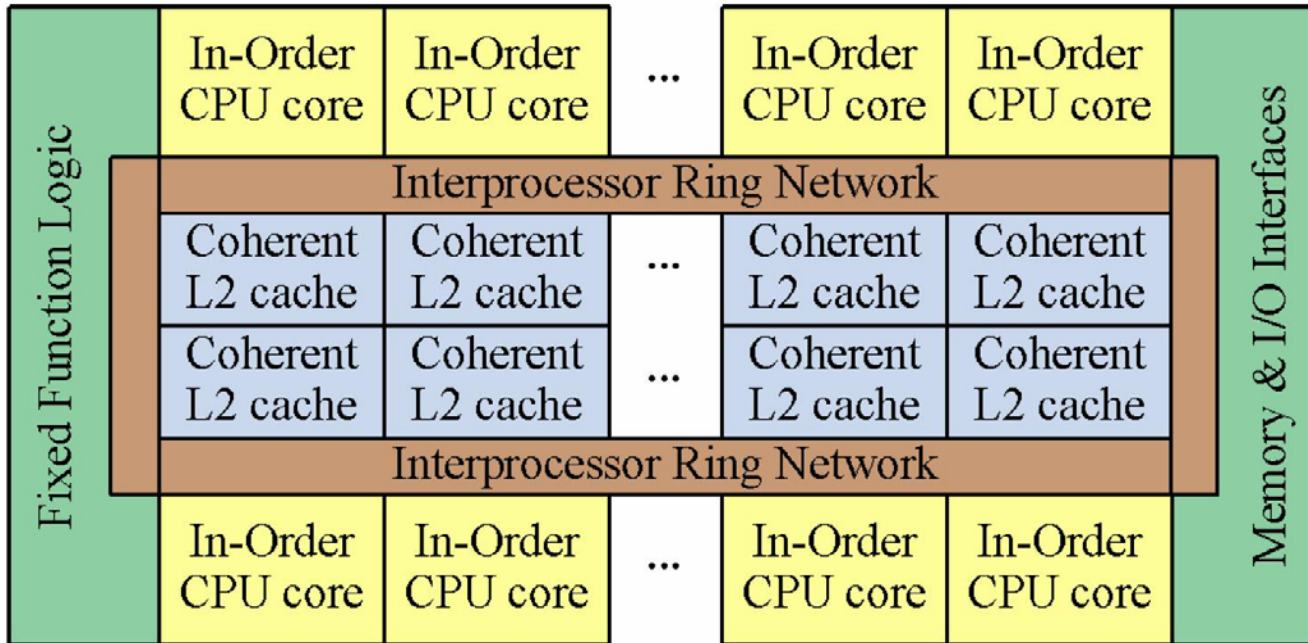
Cell Overview



- Single Power Processing Element (PPE)
- Multiple Synergistic Processing Elements (SPEs): Blade = 8, PS3 = 6, SpursEngine = 4
- SPEs are in different address spaces with limited local store (scratchpad) memory
- 4-way SIMD (vector) instructions



Larrabee Overview



- Designed as GPU (accelerator to a host core)
- Multiple cores
- Cores access a coherent cache-based memory hierarchy
- 16-way SIMD (vector) instructions
- *Image from Larrabee SigGraph 2008 paper*

Entry Method Structure

Interface File:

```
entry void entryName  
    ( ...passed parameters... );
```

Source File:

```
void ChareClass::entryName  
    ( ...passed parameters ... )  
    { ... function body ... }
```

Invocation:

```
objProxy.entryName(... passed parameters ...)
```



Basic Entry Method

```
entry void accum(int inArrayLen, float inArray[inArrayLen]);
```

```
void ChareObj::accum(int inArrayLen, float* inArray) {  
    if (inArrayLen != localArrayLen) return;  
    for (int i = 0; i < inArrayLen; ++i)  
        localArray[i] = localArray[i] + inArray[i];  
}
```

To Invoke: myChareObj.accum(someFloatArray_len,
someFloatArray_ptr);



Extensions to Charm++

SIMD Instruction Abstraction
&
Accelerated Entry Methods



SIMD Abstraction

- Abstract SIMD instructions supported by multiple architectures
 - Currently adding support for: SSE (x86), AltiVec/VMX (PowerPC; PPE), SIMD instructions on SPEs, and Larrabee
 - Generic C implementation when no direct architectural support is present
 - Types: vecf, vecdf, veci, ...
 - Operations: vaddf, vmulf, vsqrtf, ...



SIMD Instruction Abstraction

```
entry void accum(int inArrayLen, align(sizeof(vecf)) float inArray[inArrayLen]);
```

```
void ChareObj::accum(int inArrayLen, float* inArray) {  
  
    if (inArrayLen != localArrayLen) return;  
  
    vecf* inArrayVec = (vecf*)inArray;  
    vecf* localArrayVec = (vecf*)localArray;  
    int arrayVecLen = inArrayLen / vecf_numElems;  
    for (int i = 0; i < arrayVecLen; ++i)  
        localArrayVec[i] = vaddf(localArrayVec[i], inArrayVec[i]);  
  
    for (int i = arrayVecLen * vecf_numElems; i < inArrayLen; ++i)  
        localArray[i] = localArray[i] + inArray[i];  
}
```



To Invoke: myChareObj.accum(someFloatArray_len, someFloatArray_ptr);



Accelerated Entry Methods

- Targets computationally intensive code
 - Execute on accelerator if one or more supported accelerators are present
 - Otherwise, execute on host core
- Structure based on standard entry methods
 - Data dependencies expressed via messages (passed parameters)
 - Code is self-contained (accesses message and chare object data)
- The runtime system manages...
 - Data movement: DMAs automatically overlapped with work on the accelerator(s)
 - Scheduling: based on data dependencies (messages and chare objects)
- Multiple independently written portions of code share the same accelerator(s)



Accel Entry Method Structure

Basic

Interface File:

```
entry void entryName  
    ( ...passed parameters... );
```

Source File:

```
void ChareClass::entryName  
    ( ...passed parameters ... )  
    { ... function body ... }
```

VS.

Accelerated

Interface File:

```
entry [accel] void entryName  
    ( ...passed parameters... )  
    [ ...local parameters... ]  
    { ... function body ... }  
    callback_member_function;
```

Invocation (both): chareObj.entryName(... passed parameters ...)



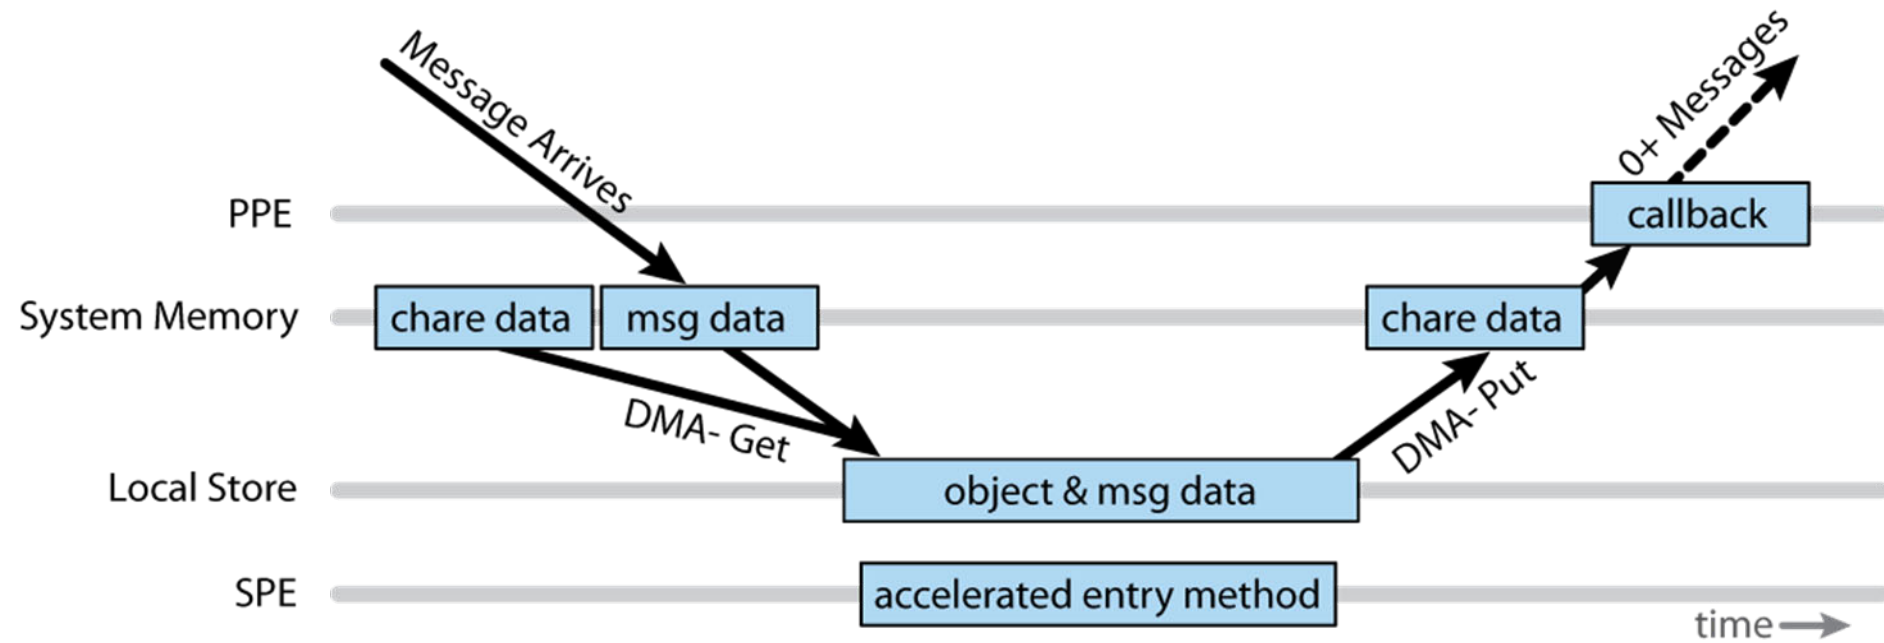
Same Entry Method Accelerated

```
entry [accel] void accum(int inArrayLen, align(sizeof(vecf)) float inArray[inArrayLen])  
    [ readOnly : int localArrayLen <impl_obj->localArrayLen>,  
      readWrite : float localArray[localArrayLen] <impl_obj->localArray> ] {  
  
    if (inArrayLen != localArrayLen) return;  
    vecf* inArrayVec = (vecf*)inArray;  
    vecf* localArrayVec = (vecf*)localArray;  
    int arrayVecLen = inArrayLen / vecf_numElems;  
  
    for (int i = 0; i < arrayVecLen; ++i)  
        localArrayVec[i] = vaddf(localArrayVec[i], inArrayVec[i]);  
  
    for (int i = arrayVecLen * vecf_numElems; i < inArrayLen; ++i)  
        localArray[i] = localArray[i] + inArray[i];  
  
} accum_callback;
```

To Invoke: myChareObj.accum(someFloatArray_len, someFloatArray_ptr);



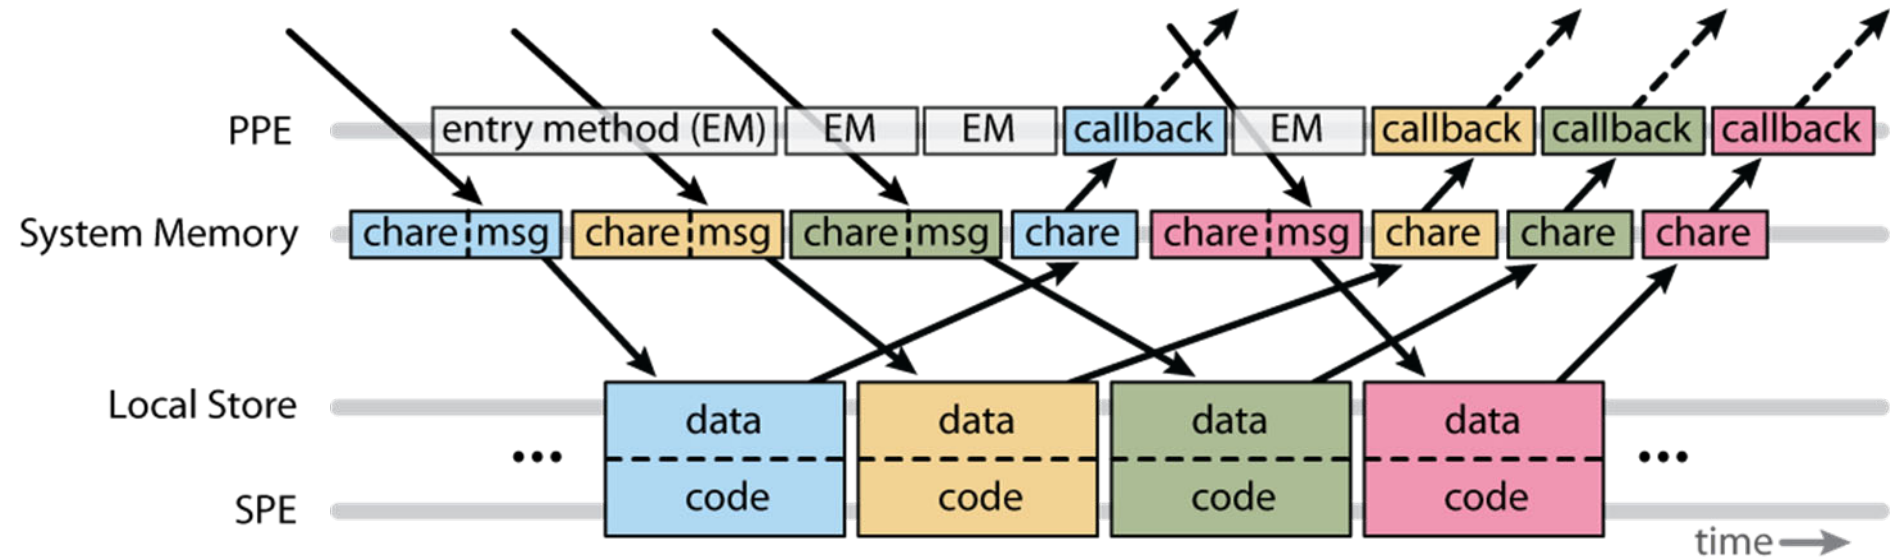
Timeline of Events



- Runtime system...
 - Directs data movement (messages & DMAs)
 - Schedules accelerated entry methods and callbacks



Communication Overlap



- Data movement automatically overlapped with accelerated entry method execution on SPEs and entry method execution on PPE

A Step Further: Heterogeneity

- Code already portable between systems with and without accelerators
 - Have runtime system account for host core differences by automatically modifying application data (message data)
 - Then able to use all cores (host and accelerators)



Summary

- Common SIMD instructions abstracted via SIMD instruction abstraction
- Accelerators available to Charm++ programs via accelerated entry methods
 - Currently targeting Cell and Larrabee
 - Ease programmer burden
- Support for heterogeneous systems
- *Disclaimer: Still under development*

