

A Multiparadigm Approach to Parallel Programming

Presented by
Aaron Becker, Pritish Jetley,
and Phil Miller

11/30/2009

Multiparadigm Computing

- A parallel computing paradigm is a way of expressing concurrency in an application
- Multiparadigm applications use multiple paradigms together
- For example, mixed-mode MPI/OpenMP applications

Some Parallel Paradigms

- MPI (Message passing)
- OpenMP (Shared memory)
- OpenCL (Accelerator)
- Charm++ (Message driven)
- MapReduce

Not Parallel Paradigms

- C++
- Object-oriented programming
- Clusters

Why Multiparadigm?

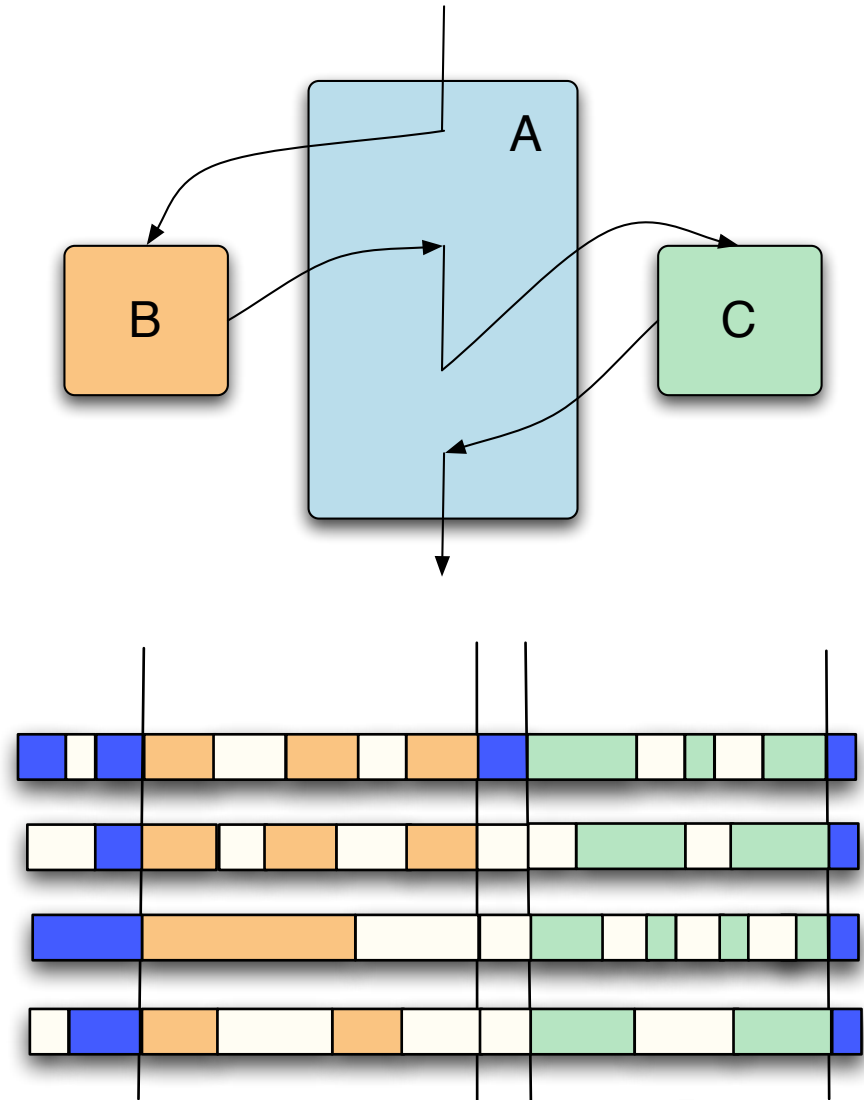
- Interoperability
- Flexibility
- Efficiency
- Simplicity

Composing Parallel Modules

- High-performance, scalable libraries are very valuable.
- What if your library
 - requires MPI?
 - assumes it will run by itself?
 - has to coexist with other parallel libraries?


Composing Parallel Modules

Sequentialization





Composing Parallel Modules

Space Division



Composing Parallel Modules

Composition
(the Charm++ model)



Benefits of Parallel Composition

- Flexibility
 - No need to statically allocate hardware resources.
 - No artificial synchronization points.
- Utilization
 - Overlap of idle time from one library with work from another.
 - Each library need not scale to the whole machine.



Simplicity

- Sometimes you don't need a fully general programming model
- Complete freedom implies the freedom to express all possible incorrect programs



Incomplete Models

Downside:

some things won't be expressible



Incomplete Models

Upside:
some things won't be expressible



Incomplete Models

Upside:

- stronger safety guarantees

- easier to optimize

- possibility for greater expressiveness



Incomplete Models

- Idea: impose restrictions on the programming model in exchange for improved safety or performance.
- Use the restricted model where it fits best, and fall back to general-purpose models elsewhere.



Incomplete Models in Charm

Multiphase Shared Arrays

disciplined shared memory

Charisma


static data-flow

Tied together with the Charm RTS



Multiphase Shared Arrays (MSA)

- An MSA application consists of:
 - Multiple computing threads
 - Data arrays (MSAs), whose elements are accessible by all threads
 - A local cache of MSA data for each thread



MSA Phases

- Observation: many shared memory applications can be split into phases.
- These phases are part of the algorithm, but aren't expressed directly in code.



MSA Phases

- In MSA, each shared array has explicit phases, e.g.
 - Read only
 - Exclusive Write
 - Accumulate
- Synchronization at phase boundaries



Safety Properties

- Provably race free
- All MSA operations obey phase semantics
- All threads accessing an MSA agree on the same sequence of phases



Asynchrony

- On remote access, fetch a remote chunk of the array into local cache
- While waiting, other work can be scheduled
 - With overdecomposition, other MSA threads can work.
 - Unrelated chares using other programming models can run, in a multiparadigm application.



Example: Plimpton MD

```
for timestep = 0 to Tmax {  
  // Phase I : Force Computation: for a section of the interaction matrix  
  for i = i_start to i_end  
    for j = j_start to j_end  
      if (nbrlist[i][j]) { // nbrlist enters ReadOnly mode  
        force = calculateForce(coords[i], atominfo[i], coords[j], atominfo[j]);  
        forces[i] += force; // Accumulate mode  
        forces[j] += -force;  
      }  
    nbrlist.sync(); forces.sync(); coords.sync();  
  
    for k = myAtomsbegin to myAtomsEnd // Phase II : Integration  
      coords[k] = integrate(atominfo[k], forces[k]); // WriteOnly mode  
    coords.sync(); atominfo.sync(); forces.sync();  
  
    if (timestep %8 == 0) { // Phase III: update neighbor list every 8 steps  
      for i = i_start to i_end  
        for j = j_start to j_end
```

